

# MongoDB to Power BI Real-time Data Pipeline

## Complete Setup & Operation Guide

---

### Prerequisites

#### Required Software (Windows 11)

##### 1. Docker Desktop for Windows

- Download: <https://www.docker.com/products/docker-desktop>
- Enable WSL 2 backend during installation

##### 2. Python 3.8 or higher

- Download: <https://www.python.org/downloads/>
- ⚠️ Check "Add Python to PATH" during installation

##### 3. Microsoft ODBC Driver 18 for SQL Server

- Download: <https://learn.microsoft.com/en-us/sql/connect/odbc/download-odbc-driver-for-sql-server>
- Choose "Download ODBC Driver 18 for SQL Server (x64)"

##### 4. Power BI Desktop

- Download: <https://powerbi.microsoft.com/desktop/>

### Required Files

All files should be in the same folder:

- requirements.txt
  - docker-compose.yml
  - Dockerfile.dockerfile
  - complete-historical-migrator.py
  - comprehensive-stream-processor.py
  - mongodb-source-config-complete.json
  - create\_tables.py
-

## Initial Setup (One-Time)

### Step 1: Install Python Dependencies

```
powershell

# Create and activate virtual environment (recommended)
python -m venv .venv
.venv\Scripts\activate

# Install all dependencies
pip install -r requirements.txt
```

### Step 2: Start Docker Infrastructure

```
bash

# Start all services (Kafka, Zookeeper, SQL Server)
docker-compose up -d

# Verify all containers are running
docker ps
```

**Expected output:** You should see 5 containers running:

- zookeeper
- kafka
- kafka-connect
- kafka-ui
- sqlserver

### Step 3: Wait for Kafka Connect to Initialize

Kafka Connect needs 2-3 minutes to start and install the MongoDB connector.

```
powershell

# Check if Kafka Connect is ready (repeat until you get a response)
Invoke-RestMethod -Uri "http://localhost:8083/connectors"

# Expected: [] (empty array means it's ready)
```

## Step 4: Create SQL Server Tables

```
bash

# Create the database tables
python create_tables.py
```

Expected output: "Tables created successfully!"

## Step 5: Load Historical Data

```
bash

# Run one-time historical data migration
python complete-historical-migrator.py
```

This will:

- Load all customers and products
- Process all historical sales and payments
- Insert aggregated data into SQL Server

Expected duration: 5-30 minutes depending on data volume

## Step 6: Register MongoDB Source Connector

```
powershell

# For Windows PowerShell
Invoke-RestMethod -Method Post -Uri "http://localhost:8083/connectors" `
  -ContentType "application/json" `
  -InFile "mongodb-source-config-complete.json"
```

Verify it's registered:

```
powershell

Invoke-RestMethod -Uri "http://localhost:8083/connectors"
# Expected: ["mongodb-source-complete"]
```

## Step 7: Start Real-time Stream Processor

```
bash
```

```
# Start the streaming processor (keep this running)
python comprehensive-stream-processor.py
```

## Expected output:

```
=====
Starting Kafka Stream Processor
=====
Loading reference data from SQL Server...
Loaded 150 customers into cache
Loaded 500 products into cache
Started thread for topic: mongodb.ClearVueDB.customer
Started thread for topic: mongodb.ClearVueDB.product
Started thread for topic: mongodb.ClearVueDB.sales
Started thread for topic: mongodb.ClearVueDB.payment
Started thread for topic: mongodb.ClearVueDB.purchase
Started thread for topic: mongodb.ClearVueDB.age_analysis
Started thread for topic: mongodb.ClearVueDB.representative
```

## Daily Operations

### Starting the Environment

#### Terminal 1: Start Docker Services

```
bash

docker-compose up -d
```

#### Terminal 2: Start Stream Processor

```
bash

# Wait 2-3 minutes after starting Docker
python comprehensive-stream-processor.py
```

That's it! The pipeline is now running and processing real-time changes.

### Stopping the Environment

#### Terminal 2: Stop Stream Processor

Press Ctrl+C

## Terminal 1: Stop Docker Services

```
bash
```

```
docker-compose down
```

## Monitoring & Verification

### Check Kafka Topics

Open browser: <http://localhost:8080>

You should see topics like:

- `mongodb.ClearVueDB.sales`
- `mongodb.ClearVueDB.customer`
- `mongodb.ClearVueDB.payment`
- etc.

### Check SQL Server Data

```
bash
```

```
# Connect to SQL Server
```

```
sqlcmd -S localhost,1433 -U sa -P akXHp5xP02YfluQ -d PowerBIDashboards
```

```
# Check record counts
```

```
SELECT 'Sales', COUNT(*) FROM HistoricalSalesSummary
```

```
UNION ALL
```

```
SELECT 'Customers', COUNT(*) FROM HistoricalCustomerPerformance
```

```
UNION ALL
```

```
SELECT 'Products', COUNT(*) FROM HistoricalProductPerformance
```

```
UNION ALL
```

```
SELECT 'Payments', COUNT(*) FROM HistoricalPaymentSummary
```

```
UNION ALL
```

```
SELECT 'Purchases', COUNT(*) FROM HistoricalPurchaseSummary
```

```
UNION ALL
```

```
SELECT 'AgeAnalysis', COUNT(*) FROM HistoricalAgeAnalysis
```

```
GO
```

## View Stream Processor Logs

The stream processor outputs logs showing:

- Records processed
  - Buffer flushes to SQL Server
  - Any errors encountered
- 



## Power BI Connection

### Connect to SQL Server

1. Open Power BI Desktop
2. Get Data → SQL Server
3. Enter connection details:
  - **Server:** localhost,1433
  - **Database:** PowerBIDashboards
  - **Authentication:** Database
  - **Username:** sa
  - **Password:** akXHrP5xP02YfluQ

## Recommended Tables & Views

### Base Tables:

- HistoricalSalesSummary
- HistoricalCustomerPerformance
- HistoricalProductPerformance
- HistoricalPaymentSummary
- HistoricalPurchaseSummary
- HistoricalAgeAnalysis

### Pre-built Views (Recommended):

- vw\_CustomerRiskAnalysis - Customer health & payment risk
- vw\_ProductProfitability - Product margins & performance
- vw\_SupplierPerformance - Supplier spending analysis

- `vw_RegionalPerformance` - Regional sales breakdown
  - `vw_CashFlowAnalysis` - Daily cash flow tracking
- 

## Testing the Pipeline

### Test Real-time Updates

#### Option 1: Insert test data in MongoDB

```
javascript

// Connect to MongoDB and insert a test sale
use ClearVueDB
db.sales.insertOne({
  "Customer_Number": "TEST001",
  "Trans_Date": "2025-09-29",
  "Doc_Num": "TEST-001",
  "Fin_Period": "202509",
  "Rep_Code": "REP01",
  "Trans_Type_Code": "SALE",
  "Sales_Line": [{
    "Inventory_Code": "PROD001",
    "Quantity": 5,
    "Unit_Sell_Price": {"$numberDecimal": "100.00"},
    "Total_Line_Price": {"$numberDecimal": "500.00"},
    "Last_Cost": {"$numberDecimal": "60.00"}
  }]
})
```

#### Option 2: Watch the stream processor logs You should see:

```
Processed sale: Customer TEST001, Amount: 500.0
Flushing buffers to SQL Server...
```

#### Option 3: Refresh Power BI Click "Refresh" in Power BI and verify the new data appears.

---

## Sharing with Your Team

### Package for Distribution

Create a project folder with:

```
mongodb-powerbi-pipeline/  
├── docker-compose.yml  
├── Dockerfile.dockerfile  
├── complete-historical-migrator.py  
├── comprehensive-stream-processor.py  
├── mongodb-source-config-complete.json  
├── create_tables.py  
├── README.md (this guide)  
└── requirements.txt
```

Create requirements.txt:

```
pymongo==4.6.0  
pyodbc==5.0.1  
kafka-python==2.0.2
```

## Team Setup Instructions

Share this guide with your team. Each team member needs to:

1. Install Docker Desktop
2. Install Python 3.8+
3. Install ODBC Driver 18
4. Clone/copy the project folder
5. Run the setup steps above

**Important:** Only ONE person should run the historical migrator. Everyone else can just start the stream processor to monitor real-time changes.

---

## Troubleshooting

### "Cannot connect to SQL Server"

```
bash
```



```
# Check if SQL Server container is running
```

```
docker ps | grep sqlserver
```

```
# Check SQL Server logs
```

```
docker logs sqlserver
```

## "Kafka Connect not responding"

```
bash
```

```
# Restart Kafka Connect
```

```
docker-compose restart kafka-connect
```

```
# Wait 2-3 minutes and check again
```

```
curl http://localhost:8083/connectors
```

## "No data in Power BI"

1. Verify historical migration completed successfully
2. Check SQL Server has data (see monitoring section)
3. Verify stream processor is running
4. Refresh Power BI connection


## Stream processor shows errors

Check the logs for specific error messages. Common issues:

- MongoDB connection failed (check credentials)
- SQL Server connection failed (check Docker is running)
- Kafka topics not found (check Kafka Connect is registered)

---

## Security Notes

 **IMPORTANT:** The current setup uses hardcoded credentials which is fine for development/testing but NOT for production.

For production deployment:

- Use environment variables for credentials
- Change default passwords

- Implement proper network security
  - Use SSL/TLS for connections
  - Restrict SQL Server access
- 



## Architecture Overview

MongoDB Atlas (Cloud)



Kafka Connect (Change Data Capture)



Kafka Topics (Streaming)



Stream Processor (Python)



SQL Server (Local)



Power BI Dashboards

### Data Flow:

1. Changes happen in MongoDB
  2. Kafka Connect captures changes
  3. Stream processor aggregates and transforms
  4. SQL Server stores analytics-ready data
  5. Power BI visualizes in real-time
- 



## Support

If team members encounter issues:

1. Check this guide first
2. Verify all prerequisites are installed
3. Check Docker containers are running
4. Review stream processor logs
5. Test SQL Server connectivity

### Logs locations:

- Stream processor: Terminal output
- Docker services: `docker logs <container-name>`
- SQL Server: `docker logs sqlserver`